

Stability Analysis Code Structure

This document describes the object-oriented architecture of the StabilityAnalysis codebase, focusing on three core classes that implement simulations of recurrent neural networks with spike-frequency adaptation (**SFA**) and short-term synaptic depression (**STD**). These three classes were used together to create Figure 2.

System Equations

The SRNN dynamics are governed by the following system of differential equations (see system_equations.md):

$$\frac{dx_i}{dt} = \frac{-x_i + u_i + \sum_{j=1}^N w_{ij} b_j r_j}{\tau_d}$$

$$r_i = \phi \left(x_i - a_{0_i} - c \sum_{k=1}^K a_{ik} \right)$$

$$\frac{da_{ik}}{dt} = \frac{-a_{ik} + r_i}{\tau_{a_k}}$$

$$\frac{db_i}{dt} = \frac{1 - b_i}{\tau_{rec}} - \frac{b_i r_i}{\tau_{rel}}$$

Symbol	Description
x_i	Dendritic potential of neuron i
r_i	Firing rate (output of activation function ϕ)
a_{ik}	SFA variable for neuron i at timescale k
b_i	STD variable (synaptic resources available)
w_{ij}	Connection weight from neuron j to neuron i
u_i	External input to neuron i
τ_d	Dendritic time constant
τ_{a_k}	SFA time constant for timescale k
τ_{rec}, τ_{rel}	STD recovery and release time constants
c	SFA coupling strength

For a complete parameter reference, see parameter_table.md.

Class: SRNNModel

File: SRNNModel.m

SRNNModel is the main simulation class. It encapsulates network parameters, ODE integration, Lyapunov exponent computation, and visualization.

Usage Pattern

```
model = SRNNModel('n', 300, 'n_a_E', 3, 'n_b_E', 1, 'level_of_chaos', 1.0);
model.build();    % Initialize W, stimulus, initial conditions
model.run();      % Integrate ODEs (and compute Lyapunov exponents if lya_method != 'none')
model.plot();     % Generate time series plots
```

Key Properties

Category	Properties	Description
Network Architecture	<code>n</code> , <code>f</code> , <code>indegree</code>	Total neurons, E fraction, expected in-degree
RMT Parameters	<code>mu_E_tilde</code> , <code>mu_I_tilde</code> , <code>sigma_E_tilde</code> , <code>sigma_I_tilde</code>	Tilde-notation stats (Harris 2023)
SFA	<code>n_a_E</code> , <code>n_a_I</code> , <code>tau_a_E</code> , <code>tau_a_I</code> , <code>c_E</code> , <code>c_I</code>	Number of SFA timescales, SFA time constant vectors (log spaced), coupling constants
STD	<code>n_b_E</code> , <code>n_b_I</code> , <code>tau_b_E_rec</code> , <code>tau_b_E_rel</code> , etc.	Enable/disable STD (0 or 1), and time constants if enabled
Dynamics	<code>tau_d</code>	Dendritic integration time constant
Nonlinearity	<code>activation_function</code> , <code>activation_function_derivative</code>	Firing rate nonlinearity and its derivative
Simulation	<code>fs</code> , <code>T_range</code> , <code>ode_solver</code> , <code>ode_opts</code>	Sampling freq, time interval, solver
Lyapunov	<code>lya_method</code> , <code>lya_T_interval</code>	Method ('benettin', 'qr', 'none'), lyapunov rescaling interval

Core Methods

Method	Description
<code>build()</code>	Computes derived params, creates <code>W</code> via <code>RMTMatrix</code> , generates stimulus, initializes state
<code>run()</code>	Integrates ODEs, computes Lyapunov exponents, decimates state for plotting
<code>plot()</code>	Generates 6-panel time series figure (input, <code>x</code> , <code>r</code> , <code>a</code> , <code>b</code> , Lyapunov)
<code>get_params()</code>	Returns a struct for compatibility with legacy functions

State Vector Organization

The state vector `S` is packed as:

```
S = [a_E(:); a_I(:); b_E(:); b_I(:); x(:)]
```

where: - `a_E`: E adaptation variables (`n_E` x `n_a_E`, flattened) - `a_I`: I adaptation variables (`n_I` x `n_a_I`, flattened) - `b_E`: E STD variables (`n_E` x `n_b_E`, flattened) - `b_I`: I STD variables (`n_I` x `n_b_I`, flattened) - `x`: Dendritic states for all neurons (`n` x 1)

Equation Implementation in `dynamics_fast()`

The static method `dynamics_fast(t, S, params)` implements the ODEs. Key steps:

1. **Interpolate external input:** `u = params.u_interpolant(t)'`
2. **Unpack state:** Extract `a_E`, `a_I`, `b_E`, `b_I`, `x` from `S`
3. **Compute effective potential:** `x_eff = x - c_E * sum(a_E, 2)` for E neurons (and similarly for I)
4. **Apply STD:** `b = ones(n,1)` with `b(E_indices) = b_E` if enabled
5. **Compute firing rate:** `r = activation_fn(x_eff)`
6. **Compute derivatives:**
 - `dx_dt = (-x + W * (b .* r) + u) / tau_d`
 - `da_E_dt = (r(E_indices) - a_E) ./ tau_a_E`
 - `db_E_dt = (1 - b_E) / tau_b_E_rec - (b_E .* r(E_indices)) / tau_b_E_rel`
7. **Pack derivatives:** `dS_dt = [da_E_dt(:); da_I_dt(:); db_E_dt(:); db_I_dt(:); dx_dt]`

Lyapunov Methods

The `lya_method` property controls how Lyapunov exponents are computed:

Method	Output	Description	Time Complexity	Memory
'benettin' (default)	LLE only	Computes local and finite-time largest Lyapunov exponent using a shadow trace (single perturbation vector rescaled periodically)	$O(T \cdot K_steps \cdot N_states)$	$O(N_states)$
'qr'	Full spectrum	Computes all N_states Lyapunov exponents via QR decomposition of the tangent space; integrates N_states^2 variational equations	$O(T \cdot K_steps \cdot N_states^3)$	$O(N_states^2 + T \cdot N_states)$
'none'	—	Skips Lyapunov computation entirely	—	—

where: - **N_states** = system dimension (**N_sys_eqs**), not to be confused with network size N - **T** = number of rescaling intervals = $(T_end - T_start) / lya_dt_interval$ - **K_steps** = average ODE substeps per rescaling interval (depends on solver tolerances and dynamics stiffness)

Recommendation: Use 'benettin' for large networks ($N_states > 100$) or parameter sweeps. Use 'qr' only when the full spectrum is needed and N_states is small (the N_states^3 scaling makes it prohibitive for $N_states > \sim 200$).

Class: RMTMatrix

File: RMTMatrix.m

RMTMatrix constructs sparse connectivity matrices following Random Matrix Theory (Harris et al., 2023). It enforces Dale's law by separating excitatory and inhibitory populations with distinct mean weights.

Usage Pattern

```
rmt = RMTMatrix(N);
rmt.alpha = indegree / N;      % Sparsity
rmt.f = 0.5;                   % E fraction
rmt.mu_tilde_e = 3 * D;        % E mean (tilde notation)
rmt.mu_tilde_i = -4 * D;       % I mean (tilde notation)
rmt.sigma_tilde_e = D;         % E std dev
rmt.sigma_tilde_i = D;         % I std dev
W = rmt.W;                     % Access triggers construction
```

Key Properties

Property	Description
<code>N</code>	Network size
<code>alpha</code>	Connection probability (sparsity)
<code>f</code>	Fraction of excitatory neurons
<code>mu_tilde_e, mu_tilde_i</code>	Normalized population means
<code>sigma_tilde_e, sigma_tilde_i</code>	Normalized population std devs
<code>zrs_mode</code>	Zero row sum mode: 'none', 'ZRS', 'SZRS', 'Partial-SZRS'

Dependent Properties (Computed)

Property	Formula	Description
mu_se	$\alpha \cdot \tilde{\mu}_E$	Sparse excitatory mean
mu_si	$\alpha \cdot \tilde{\mu}_I$	Sparse inhibitory mean
R	$\sqrt{N(f\sigma_{se}^2 + (1-f)\sigma_{si}^2)}$	Theoretical spectral radius (Eq 18)
lambda_0	$N(f\mu_{se} + (1-f)\mu_{si})$	Outlier eigenvalue (Eq 17)

W Matrix Construction

The weight matrix W is constructed as:

```
W_dense = (A * D) + M;    % A: Gaussian random, D: variance diagonal, M: low-rank mean
W = S .* W_dense;        % S: sparse binary mask
```

where: - A is an N x N Gaussian random matrix (mean 0, variance 1) - D = diag([sigma_tilde_e, ..., sigma_tilde_e, sigma_tilde_i, ..., sigma_tilde_i]) encodes population variance - M = u * v' is the rank-1 mean structure with v = [mu_tilde_e, ..., mu_tilde_e, mu_tilde_i, ..., mu_tilde_i]' - S is a Bernoulli (0 or 1) sparsity mask with connection probability alpha

Class: ParamSpaceAnalysis

File: ParamSpaceAnalysis.m

ParamSpaceAnalysis performs multi-dimensional gridded parameter space analysis of stability (and other metrics). The user can choose any combination of model parameters to vary. By default it simulates all parameter combinations across four conditions: no adaptation, SFA only, STD only, and SFA + STD. This enables a systematic comparison of network behavior under different parameter ranges and conditions. The parallel computing toolbox is highly recommended but not required.

Usage Pattern

```
psa = ParamSpaceAnalysis('n_levels', 5, 'note', 'f_sweep');
psa.add_grid_parameter('f', [0.4, 0.6]);    % Sweep E fraction (range mode)
psa.add_grid_parameter('reps', [1:10]);      % 10 repetitions per point (explicit mode), otherwise default
psa.model_defaults.n = 300;                  % Set constant parameters
psa.run();                                    % Execute analysis
psa.plot('metric', 'LLE');                    % Visualize results
```

add_grid_parameter Behavior

The add_grid_parameter(param_name, values) method supports two modes:

Mode	Input	Behavior
Range	1x2 vector [min, max]	Evenly divides the range into n_levels values using linspace(min, max, n_levels)
Explicit	Vector with 3+ elements	Uses the exact values provided (ignores n_levels for this parameter)

Examples:

```
psa.add_grid_parameter('f', [0.4, 0.6]);    % Range mode: 5 levels + [0.4, 0.45, 0.5, 0.55, 0.6]
psa.add_grid_parameter('reps', [1, 2, 3, 4, 5]); % Explicit mode: uses [1, 2, 3, 4, 5] directly
```

Note: Each call to `add_grid_parameter` adds a dimension to the parameter space. The total number of simulations grows **multiplicatively**:

Total = `n_conditions` x `n_reps` x (`n_levels_1` x `n_levels_2` x ...)

For example, 3 grid parameters with 5 levels each, 10 repetitions, and 4 conditions yields: $4 \times 10 \times 5^3 = 5,000$ total simulations.

Key Features

- **Multi-dimensional grid:** All combinations of specified parameters are tested
- **Four adaptation conditions** (default): | Condition | `n_a_E` | `n_b_E` | Description | |-----|-----|-----|
|-----| | `no_adaptation` | 0 | 0 | Baseline | | `sfa_only` | 3 | 0 | SFA enabled | | `std_only` | 0 | 1 | STD enabled | | `sfa_and_std` | 3 | 1 | Both enabled |
- **Same W across conditions:** Each grid point uses identical connectivity for fair comparison
- **Batched parallel execution:** Uses `parfor` with checkpointing for resume capability
- **Randomized order:** Allows representative early stopping

Output Structure

Results are saved to `data/param_space/<timestamped_folder>/`:

```
param_space_<note>_nLevs_<N>_<timestamp>/
|-- param_space_summary.mat          # Grid configuration
|-- psa_object.mat                  # Serialized PSA object
|-- no_adaptation/
|   +-- param_space_results_no_adaptation.mat
|-- sfa_only/
|   +-- param_space_results_sfa_only.mat
|-- std_only/
|   +-- param_space_results_std_only.mat
+-- sfa_and_std/
    +-- param_space_results_sfa_and_std.mat
```

Each result struct contains: `LLE`, `mean_rate`, `mean_synaptic_output`, `config`, `success`, and optionally `local_lya` time series.

Supporting Functions

The `src/` directory contains additional utilities:

Category	Files	Description
Lyapunov	<code>compute_lyapunov_exponent.m</code> , <code>lyapunov_spectrum_qr.m</code> , <code>lyapunov_benettin.m</code>	LLE, and full spectrum computation
Jacobian	<code>compute_Jacobian_fast.m</code> , <code>compute_Jacobian_at_indices.m</code>	Sparse Jacobian construction
Stimulus	<code>generate_external_input.m</code>	Random step-function input generation
Activation	<code>piecewiseSigmoid.m</code> , <code>piecewiseSigmoidDerivative.m</code>	Nonlinearity and its derivative
Plotting	<code>plot_SRNN_tseries.m</code> , <code>plot_SRNN_combined_tseries.m</code> , etc.	Visualization utilities
Analysis	<code>load_and_make_unit_histogram.m</code> , <code>load_and_plot_lle_by_stim_period.m</code>	Post-hoc analysis of PSA results

See Also

- `parameter_table.md` – Complete parameter reference
- `system_equations.md` – Mathematical model
- `Script_Notes.md` – How to run the Figure 2 scripts